

Training reward curves, mean $\pm$ std across seeds (no smoothing)

