

Training reward curves, mean  $\pm$  std across seeds (no smoothing)

